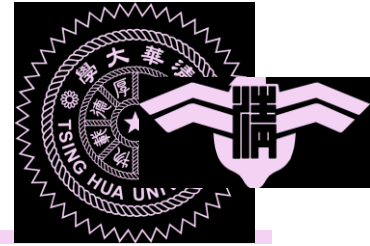


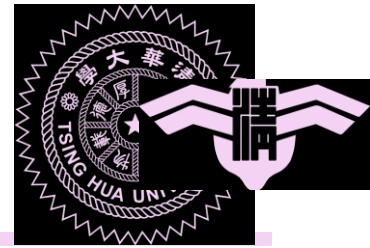
Cache Replacement Policies and Signature-based Hit Prediction (SHiP) Cache

Memory Systems HW1



Outline

- Cache replacement policies
- Simulator introduction
- NTHU Online Judge introduction



Replacement Policy

- Each set of an N-way cache has N blocks
- Upon a miss, the cache can consider giving up one cached block to store the requested block
- A selection policy is needed
 - The selection affects performance (not correctness)
- Some cache replacement policies
 1. LRU (least recently used)
 2. NRU (not recently used)
 3. SRRIP (static re-reference interval prediction)
 4. SHiP+SRRIP

1. LRU

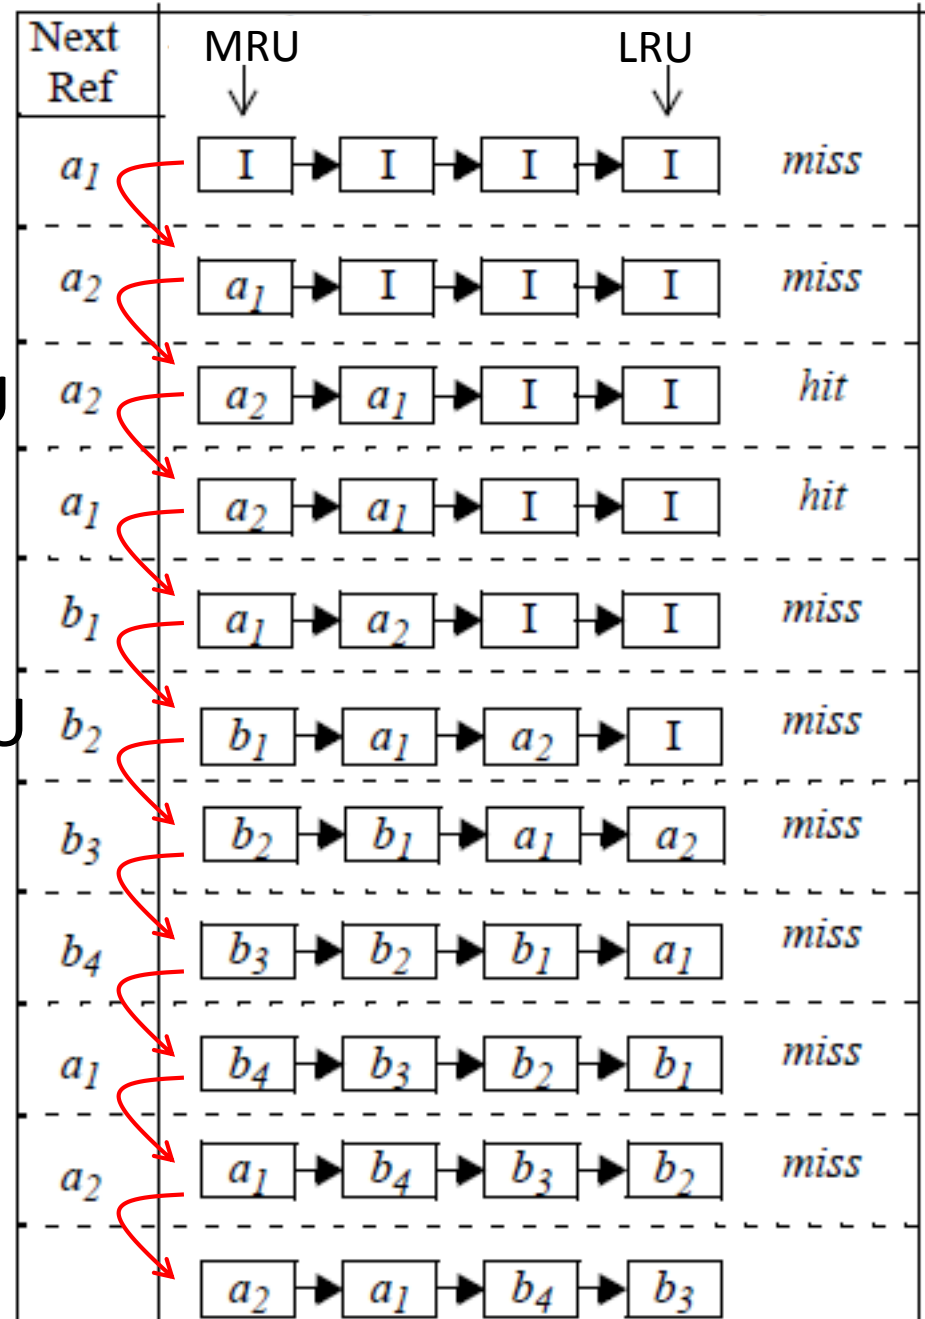
Cache Hit:

(i) move the block to MRU

Cache Miss:

(i) replace the LRU block

(ii) move the block to MRU



2. NRU

Cache Hit:

(i) set the nru-bit of the block to '0'

Cache Miss:

(i) search for the first '1' from the left

(ii) if '1' is found go to step (v)

(iii) set all nru-bits to '1'

(iv) goto step (i)

(v) replace the block and set the nru-bit to '1'

Next Ref	nru-bit				
a_1	I_1	I_1	I_1	I_1	miss
a_2	a_1_0	I_1	I_1	I_1	miss
a_2	a_1_0	a_2_0	I_1	I_1	hit
a_1	a_1_0	a_2_0	I_1	I_1	hit
b_1	a_1_0	a_2_0	I_1	I_1	miss
b_2	a_1_0	a_2_0	b_1_0	I_1	miss
b_3	a_1_0	a_2_0	b_1_0	b_2_0	miss
b_4	b_3_0	a_2_1	b_1_1	b_2_1	miss
a_1	b_3_0	b_4_0	b_1_1	b_2_1	miss
a_2	b_3_0	b_4_0	a_1_0	b_2_1	miss
	b_3_0	b_4_0	a_1_0	a_2_0	

3. SRRIP

Cache Hit:

(i) set the RRPV of the block to '0'

Cache Miss:

(i) search for the first '3' from the left


(ii) if '3' is found go to step (v)

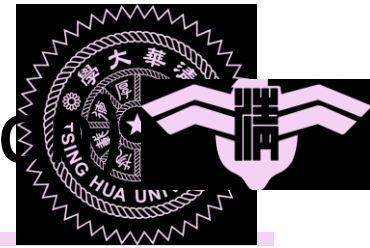
(iii) increment all RRPVs

(iv) goto step (i)

(v) replace the block and set the RRPV to '2'

Next Ref	RRPV				
a_1	I_3	I_3	I_3	I_3	<i>miss</i>
a_2	a_1_2	I_3	I_3	I_3	<i>miss</i>
a_2	a_1_2	a_2_2	I_3	I_3	<i>hit</i>
a_1	a_1_2	a_2_0	I_3	I_3	<i>hit</i>
b_1	a_1_0	a_2_0	I_3	I_3	<i>miss</i>
b_2	a_1_0	a_2_0	b_1_2	I_3	<i>miss</i>
b_3	a_1_0	a_2_0	b_1_2	b_2_2	<i>miss</i>
b_4	a_1_1	a_2_1	b_3_2	b_2_3	<i>miss</i>
a_1	a_1_1	a_2_1	b_3_2	b_4_2	<i>hit</i>
a_2	a_1_0	a_2_1	b_3_2	b_4_2	<i>hit</i>
	a_1_0	a_2_0	b_3_2	b_4_2	

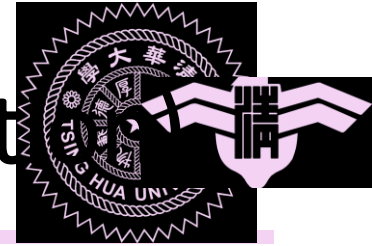
LRU (least recently used)	NRU (not recently used)	RRIP 
<div>Next Ref</div> <div>RRIP head</div> <div>RRIP tail</div>		
<div>a_1</div> <div>I → I → I → I</div> <div>miss</div>	<div>I₁ I₁ I₁ I₁</div> <div>miss</div>	<div>I₃ I₃ I₃ I₃</div> <div>miss</div>
<div>a_2</div> <div>a_1 → I → I → I</div> <div>miss</div>	<div>$a_1$₀ I₁ I₁ I₁</div> <div>miss</div>	<div>$a_1$₂ I₃ I₃ I₃</div> <div>miss</div>
<div>a_2</div> <div>a_2 → a_1 → I → I</div> <div>hit</div>	<div>$a_1$₀ $a_2$₀ I₁ I₁</div> <div>hit</div>	<div>$a_1$₂ $a_2$₂ I₃ I₃</div> <div>hit</div>
<div>a_1</div> <div>a_2 → a_1 → I → I</div> <div>hit</div>	<div>$a_1$₀ $a_2$₀ I₁ I₁</div> <div>hit</div>	<div>$a_1$₂ $a_2$₀ I₃ I₃</div> <div>hit</div>
<div>b_1</div> <div>a_1 → a_2 → I → I</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ I₁ I₁</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ I₃ I₃</div> <div>miss</div>
<div>b_2</div> <div>b_1 → a_1 → a_2 → I</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ $b_1$₀ I₁</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ $b_1$₂ I₃</div> <div>miss</div>
<div>b_3</div> <div>b_2 → b_1 → a_1 → a_2</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ $b_1$₀ $b_2$₀</div> <div>miss</div>	<div>$a_1$₀ $a_2$₀ $b_1$₂ $b_2$₂</div> <div>miss</div>
<div>b_4</div> <div>b_3 → b_2 → b_1 → a_1</div> <div>miss</div>	<div>$b_3$₀ $a_2$₁ $b_1$₁ $b_2$₁</div> <div>miss</div>	<div>$a_1$₁ $a_2$₁ $b_3$₂ $b_2$₃</div> <div>miss</div>
<div>a_1</div> <div>b_4 → b_3 → b_2 → b_1</div> <div>miss</div>	<div>$b_3$₀ $b_4$₀ $b_1$₁ $b_2$₁</div> <div>miss</div>	<div>$a_1$₁ $a_2$₁ $b_3$₂ $b_4$₂</div> <div>hit</div>
<div>a_2</div> <div>a_1 → b_4 → b_3 → b_2</div> <div>miss</div>	<div>$b_3$₀ $b_4$₀ $a_1$₀ $b_2$₁</div> <div>miss</div>	<div>$a_1$₀ $a_2$₁ $b_3$₂ $b_4$₂</div> <div>hit</div>
<div>a_2</div> <div>a_2 → a_1 → b_4 → b_3</div> <div></div>	<div>$b_3$₀ $b_4$₀ $a_1$₀ $a_2$₀</div> <div></div>	<div>$a_1$₀ $a_2$₀ $b_3$₂ $b_4$₂</div> <div></div>
<div>Cache Hit:</div> <div>(i) move block to MRU</div> <div>Cache Miss:</div> <div>(i) replace LRU block</div> <div>(ii) move block to MRU</div>	<div>Cache Hit:</div> <div>(i) set nru-bit of block to '0'</div> <div>Cache Miss:</div> <div>(i) search for first '1' from the left</div> <div>(ii) if '1' is found go to step (v)</div> <div>(iii) set all nru-bits to '1'</div> <div>(iv) goto step (i)</div> <div>(v) replace block and set nru-bit to '1'</div>	<div>Cache Hit:</div> <div>(i) set RRPV of block to '0'</div> <div>Cache Miss:</div> <div>(i) search for first '3' from the left</div> <div>(ii) if '3' is found go to step (v)</div> <div>(iii) increment all RRPVs</div> <div>(iv) goto step (i)</div> <div>(v) replace block and set RRPV to '2'</div>



4. SHiP (Signature-based Hit Prediction)

- SHiP can improve over SRRIP

SRRIP	SRRIP + SHiP
Cache Hit: (i) set RRPV of block to '0' Cache Miss: (i) search for first '3' from left (ii) if '3' found go to step (v) (iii) increment all RRPVs (iv) goto step (i) (v) replace block and set RRPV to '2'	Cache Hit: (i) set RRPV of block to '0' Cache Miss: (i) search for first '3' from left (ii) if '3' found go to step (v) (iii) increment all RRPVs (iv) goto step (i) (v) replace block and set RRPV according to SHiP



SHiP (Signature-based Hit Predict

- Some blocks have a very long reuse distance
 - Not very worth being cached
- RRPV of this kind of blocks would be better set to '3' instead of '2'
- This kind of blocks can be identified using signature
 - Memory address
 - Program counter of the load/store instruction
 - etc.

SRRIP + SHiP

Cache Hit:

(i) set RRPV of block to '0'

Cache Miss:

(i) search for first '3' from left

(ii) if '3' found go to step (v)

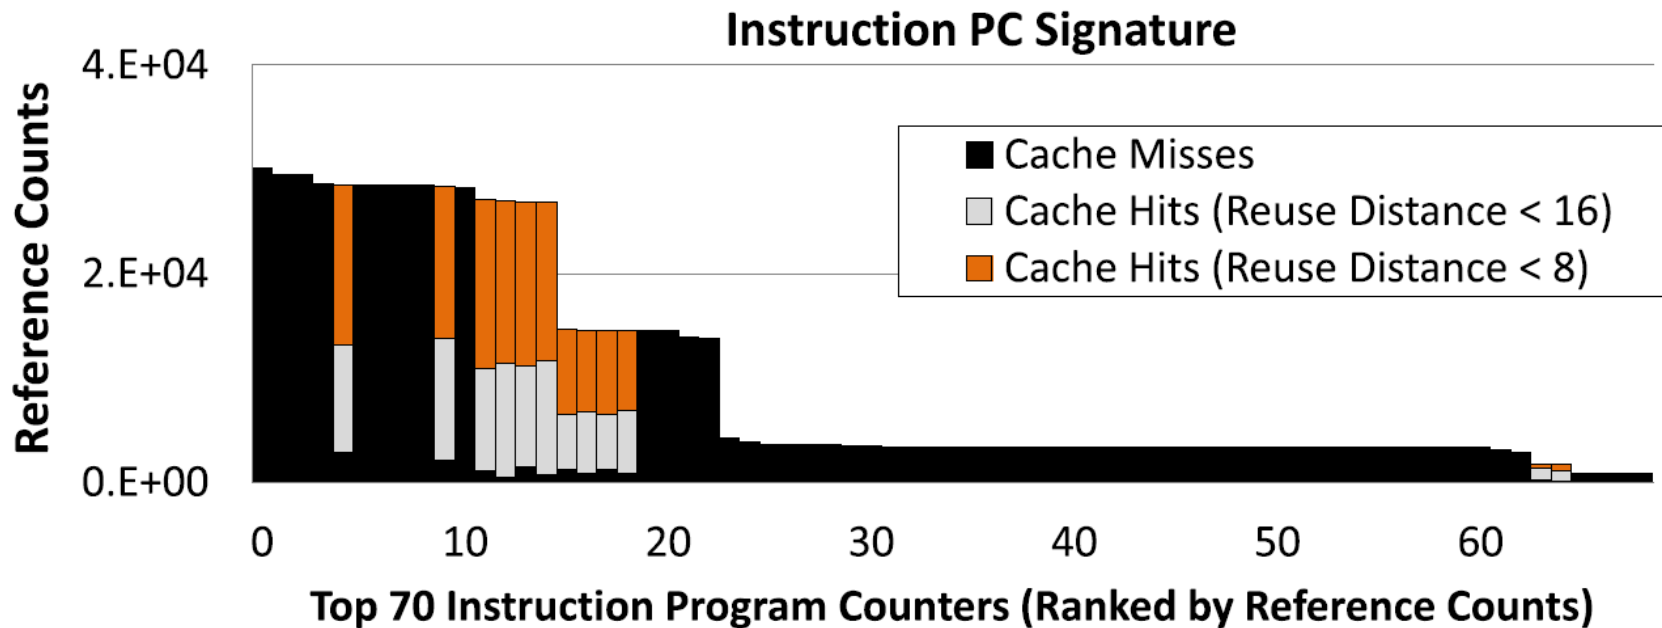
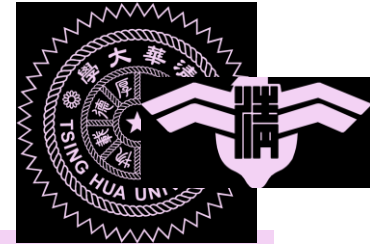
(iii) increment all RRPVs

(iv) goto step (i)

(v) replace block and set RRPV

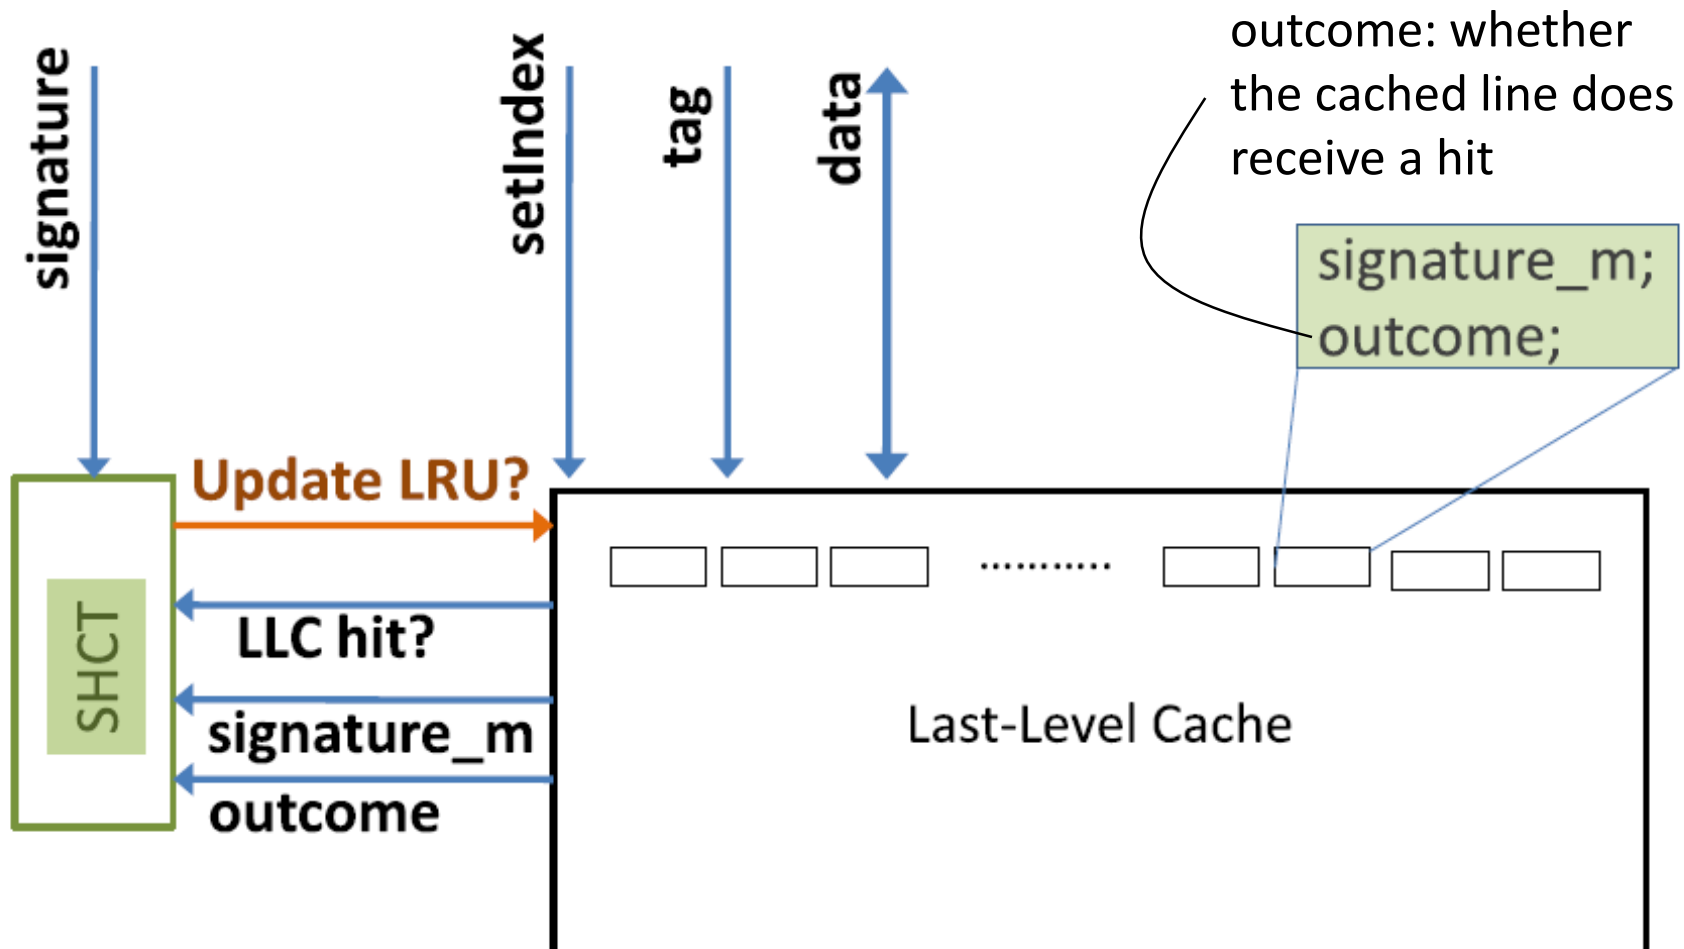
according to SHiP

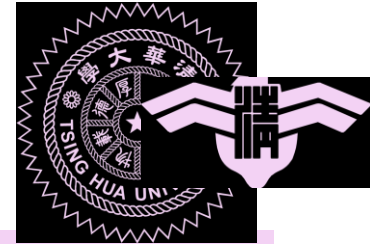
Program Counter Signature





SHiP Architecture





SHiP-PC + SRRIP

Cache Hit:

```
cache_line.outcome = true;  
Increment SHCT[signature_line];  
RRPV = 0; // same as RRIP
```

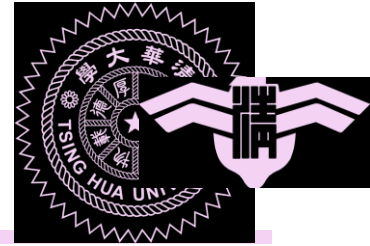
Cache Miss:

```
find an old line to evict  
if (evicted_cache_line.outcome == false)  
    Decrement SHCT[signature_line];  
insert the new line  
cache_line.outcome = false;  
cache_line.signature_line = hash(PC);  
if (SHCT[hash(PC)] == 0)  
    RRPV = 3; // not worth being cached  
else  
    RRPV = 2; // same as RRIP
```



Simulator

- SHiP+SRRIP looks good
- We need to quantify the improvement of SHiP+SRRIP over baseline design such as LRU
- Thus, we need simulator to answer the question



Simulator

- Use a Linux workstation (e.g., ws31 of NTHUEE) or use a Linux virtual machine
- Download the simulator kit from the course website
- Unpack the tarball

```
tar -xzf HW2.tgz  
cd CRC
```

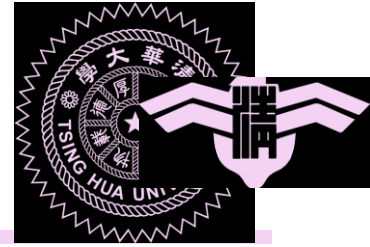
- Set the variable 'PIN_DIR' in pinkit/CONFIG/makefile.gnu.config to the absolute path for pinkit/pin-2.7-31933-gcc.3.4.6-ia32_intel64-linux/ in the kit. When you edit the makefile.gnu.config file, modify the fourth line to:

```
PIN_DIR ?= $FULL_PATH_TO_CRC_KIT/CRC/pinkit/pin-2.7-31933-  
gcc.3.4.6-ia32_intel64-linux
```



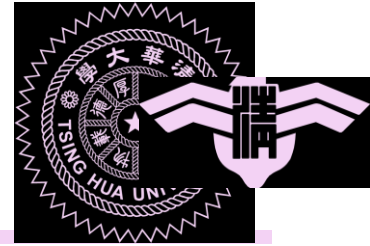
Simulator

- To compile the simulator, make sure you are in the CRC/ directory and that your gcc compiler path is set up correctly, then type (depending on whether your machine is 32 or 64 bits):
make CMPsim32
or
make CMPsim64
 - type "uname -a" in Linux to determine whether your machine is 32-bit (IA-32, i686) or 64-bit (x86_64)
- Traces are available in the CRC/traces director



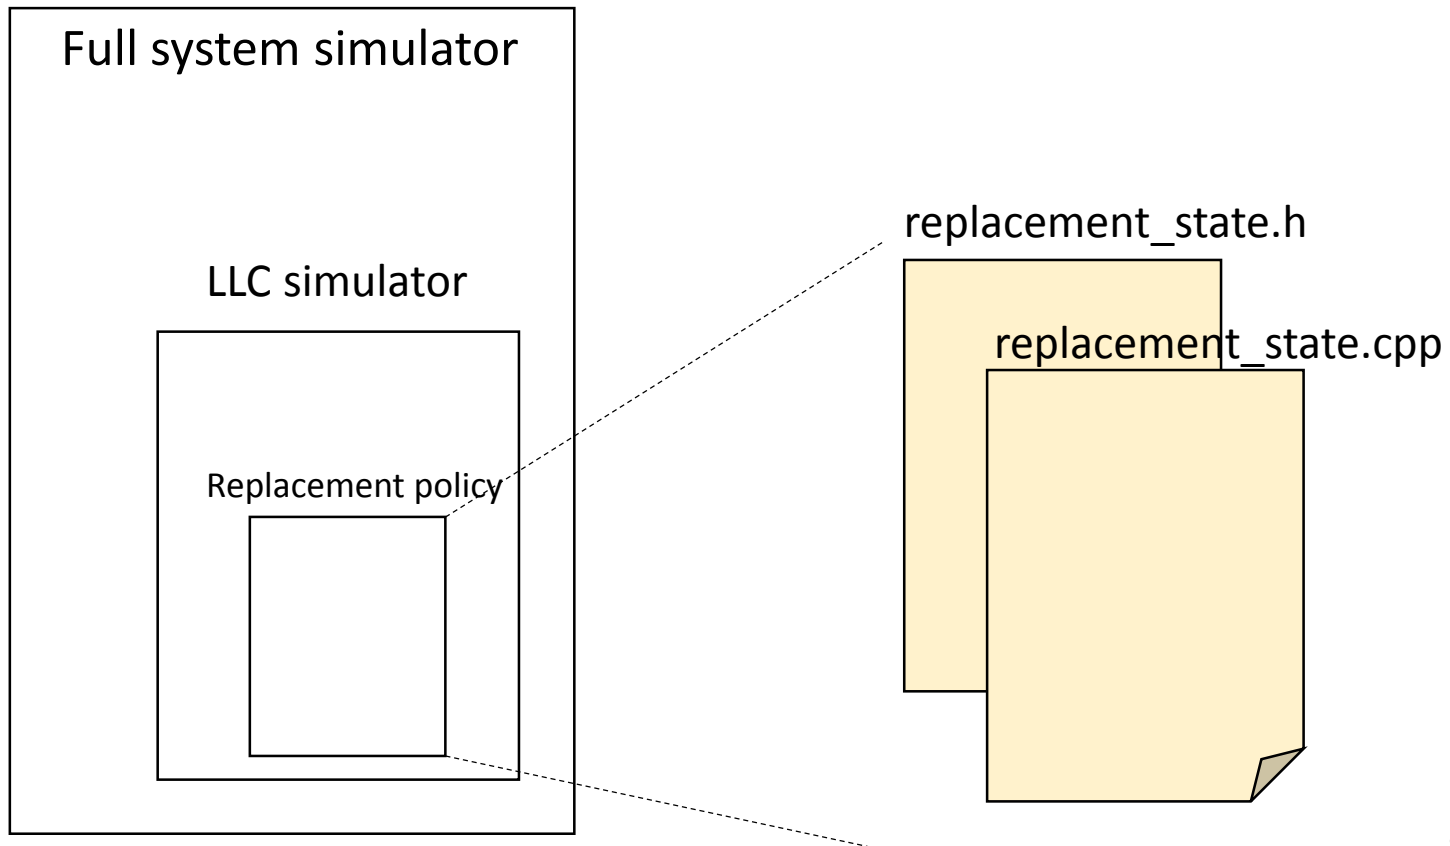
Run Simulation

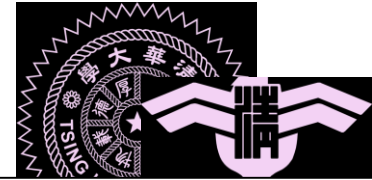
- Edit the full path information of trace1 to trace4 in the CRC/traces/ directory
 - **FULL_PATH**/CRC/traces/art.out.trace.gz
- Edit the 8th line of runall.sh in the CRC/runs/ directory
 - Change "CMPsim.usetrace.**64**" to "CMPsim.usetrace.**32**" if your machine is 32-bit



Modify Simulator

- You are allowed to modified **only** two files
 - CRC/src/LLCsim/replacement_state.h
 - CRC/src/LLCsim/replacement_state.cpp





replacement_state.h

```
class CACHE_REPLACEMENT_STATE
{
private:
    UINT32 numsets;
    UINT32 assoc;
    UINT32 replPolicy;
    LINE_REPLACEMENT_STATE **repl;
    COUNTER mytimer; // tracks # of references to the cache
    // Add extra state for cache here

public:

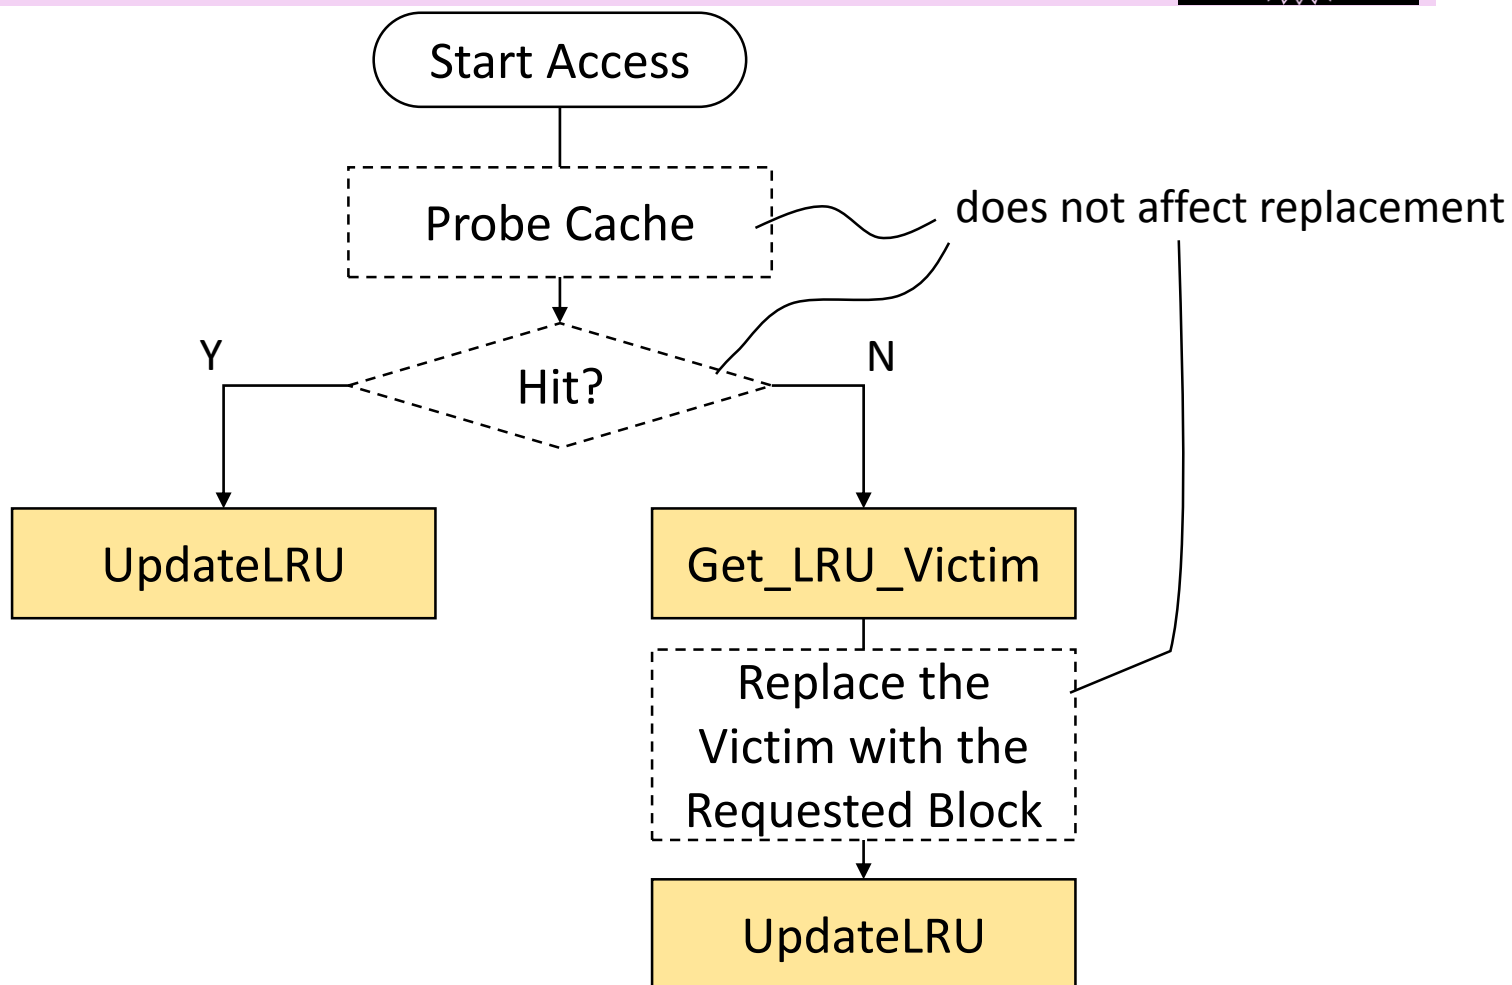
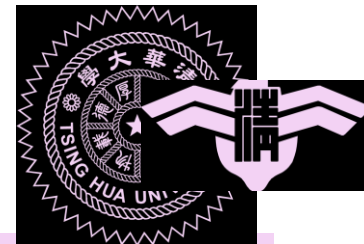
    CACHE_REPLACEMENT_STATE( UINT32 _sets, UINT32 _assoc, UINT32 _pol );
    INT32 GetVictimInSet( ... );
    void UpdateReplacementState( UINT32 setIndex, INT32 updateWayID );
    void SetReplacementPolicy( UINT32 _pol ) { replPolicy = _pol; }
    void IncrementTimer() { mytimer++; }
    void UpdateReplacementState( ... );
    ostream& PrintStats( ostream &out);

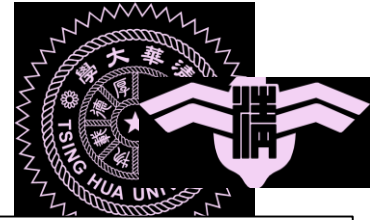
private:

    void InitReplacementState();
    INT32 Get_Random_Victim( UINT32 setIndex );
    INT32 Get_LRU_Victim( UINT32 setIndex );
    void UpdateLRU( UINT32 setIndex, INT32 updateWayID );
};
```

```
typedef struct
{
    UINT32 LRUstackposition;
    // Add extra "per-line" state here
} LINE_REPLACEMENT_STATE;
```

LRU Example





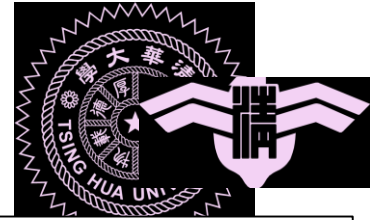
replacement_state.cpp

```
CACHE_REPLACEMENT_STATE::CACHE_REPLACEMENT_STATE( UINT32 _sets,
UINT32 _assoc, UINT32 _pol )
{

    numsets      = _sets;
    assoc        = _assoc;
    replPolicy    = _pol;

    mytimer      = 0;

    InitReplacementState();
}
```



replacement_state.cpp

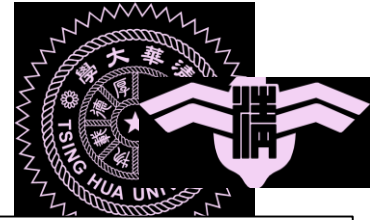
```
void CACHE_REPLACEMENT_STATE::InitReplacementState()
{
    // Create the state for sets, then create the state for the ways
    repl = new LINE_REPLACEMENT_STATE* [ numsets ];

    // ensure that we were able to create replacement state
    assert(repl);

    // Create the state for the sets
    for(UINT32 setIndex=0; setIndex<numsets; setIndex++)
    {
        repl[ setIndex ] = new LINE_REPLACEMENT_STATE[ assoc ];

        for(UINT32 way=0; way<assoc; way++)
        {
            // initialize stack position (for true LRU)
            repl[ setIndex ][ way ].LRUstackposition = way;
        }
    }

    // Contestants:  ADD INITIALIZATION FOR YOUR HARDWARE HERE
}
```

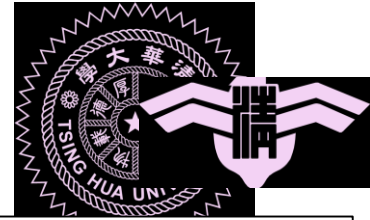


replacement_state.cpp

```
INT32 CACHE_REPLACEMENT_STATE::GetVictimInSet( UINT32 tid, UINT32 setIndex, const
LINE_STATE *vicSet, UINT32 assoc, Addr_t PC, Addr_t paddr, UINT32 accessType )
{
    // If no invalid lines, then replace based on replacement policy
    if( replPolicy == CRC_REPL_LRU )
    {
        return Get_LRU_Victim( setIndex );
    }
    else if( replPolicy == CRC_REPL_RANDOM )
    {
        return Get_Random_Victim( setIndex );
    }
    else if( replPolicy == CRC_REPL_CONTESTANT )
    {
        // ADD YOUR VICTIM SELECTION FUNCTION HERE
    }

    // We should never get here
    assert(0);

    return -1; // Returning -1 bypasses the LLC
}
```



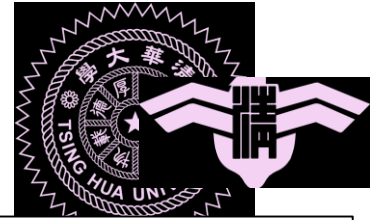
replacement_state.cpp

```
INT32 CACHE_REPLACEMENT_STATE::Get_LRU_Victim( UINT32 setIndex )
{
    // Get pointer to replacement state of current set
    LINE_REPLACEMENT_STATE *replSet = repl[ setIndex ];

    INT32 lruWay = 0;

    // Search for victim whose stack position is assoc-1
    for(UINT32 way=0; way<assoc; way++)
    {
        if( replSet[way].LRUstackposition == (assoc-1) )
        {
            lruWay = way;
            break;
        }
    }

    // return lru way
    return lruWay;
}
```

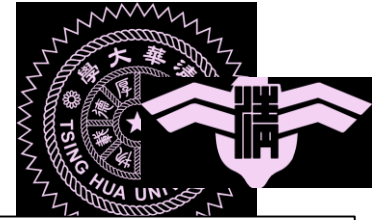


replacement_state.cpp

```
void CACHE_REPLACEMENT_STATE::UpdateReplacementState(
    UINT32 setIndex, INT32 updateWayID, const LINE_STATE *currLine,
    UINT32 tid, Addr_t PC, UINT32 accessType, bool cacheHit )
{

    if( replPolicy == CRC_REPL_LRU )
    {
        UpdateLRU( setIndex, updateWayID );
    }
    else if( replPolicy == CRC_REPL_RANDOM )
    {
        // Random replacement requires no replacement state update
    }
    else if( replPolicy == CRC_REPL_CONTESTANT )
    {
        // ADD YOUR UPDATE REPLACEMENT STATE FUNCTION HERE
        // Feel free to use any of the input parameters to make
        // updates to your replacement policy
    }

}
```

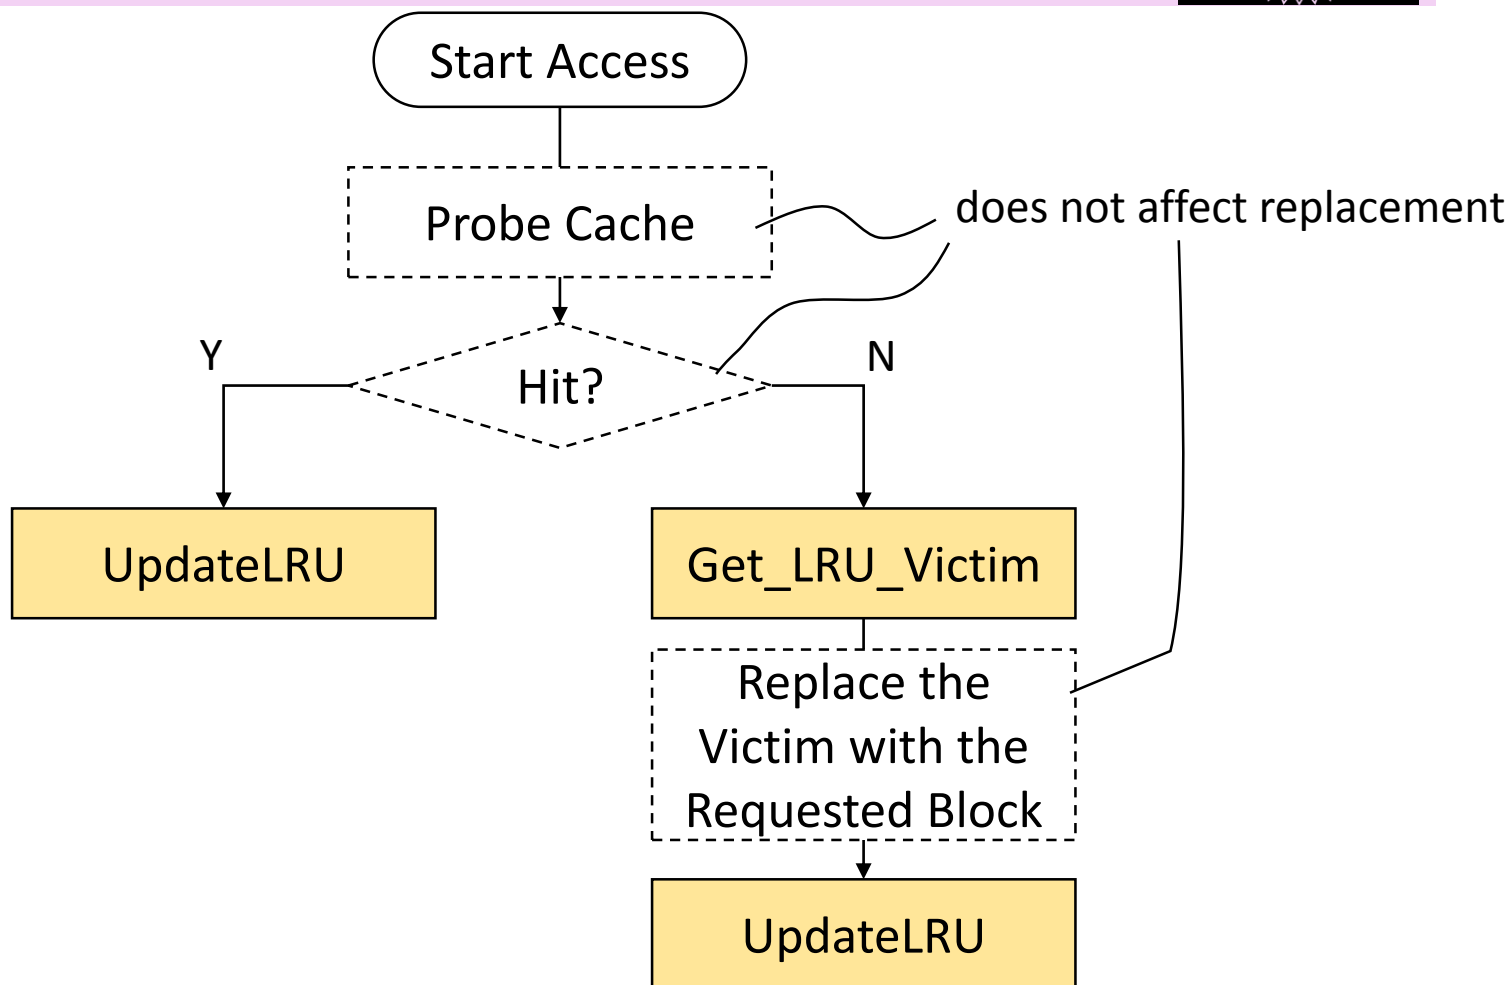
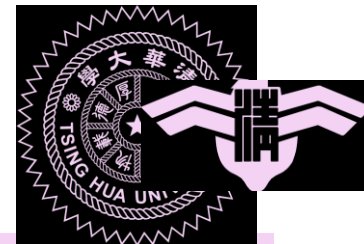
replacement_state.cpp

```
void CACHE_REPLACEMENT_STATE::UpdateLRU( UINT32 setIndex, INT32 updateWayID )
{
    // Determine current LRU stack position
    UINT32 currLRUstackposition = repl[ setIndex ][ updateWayID ].LRUstackposition;

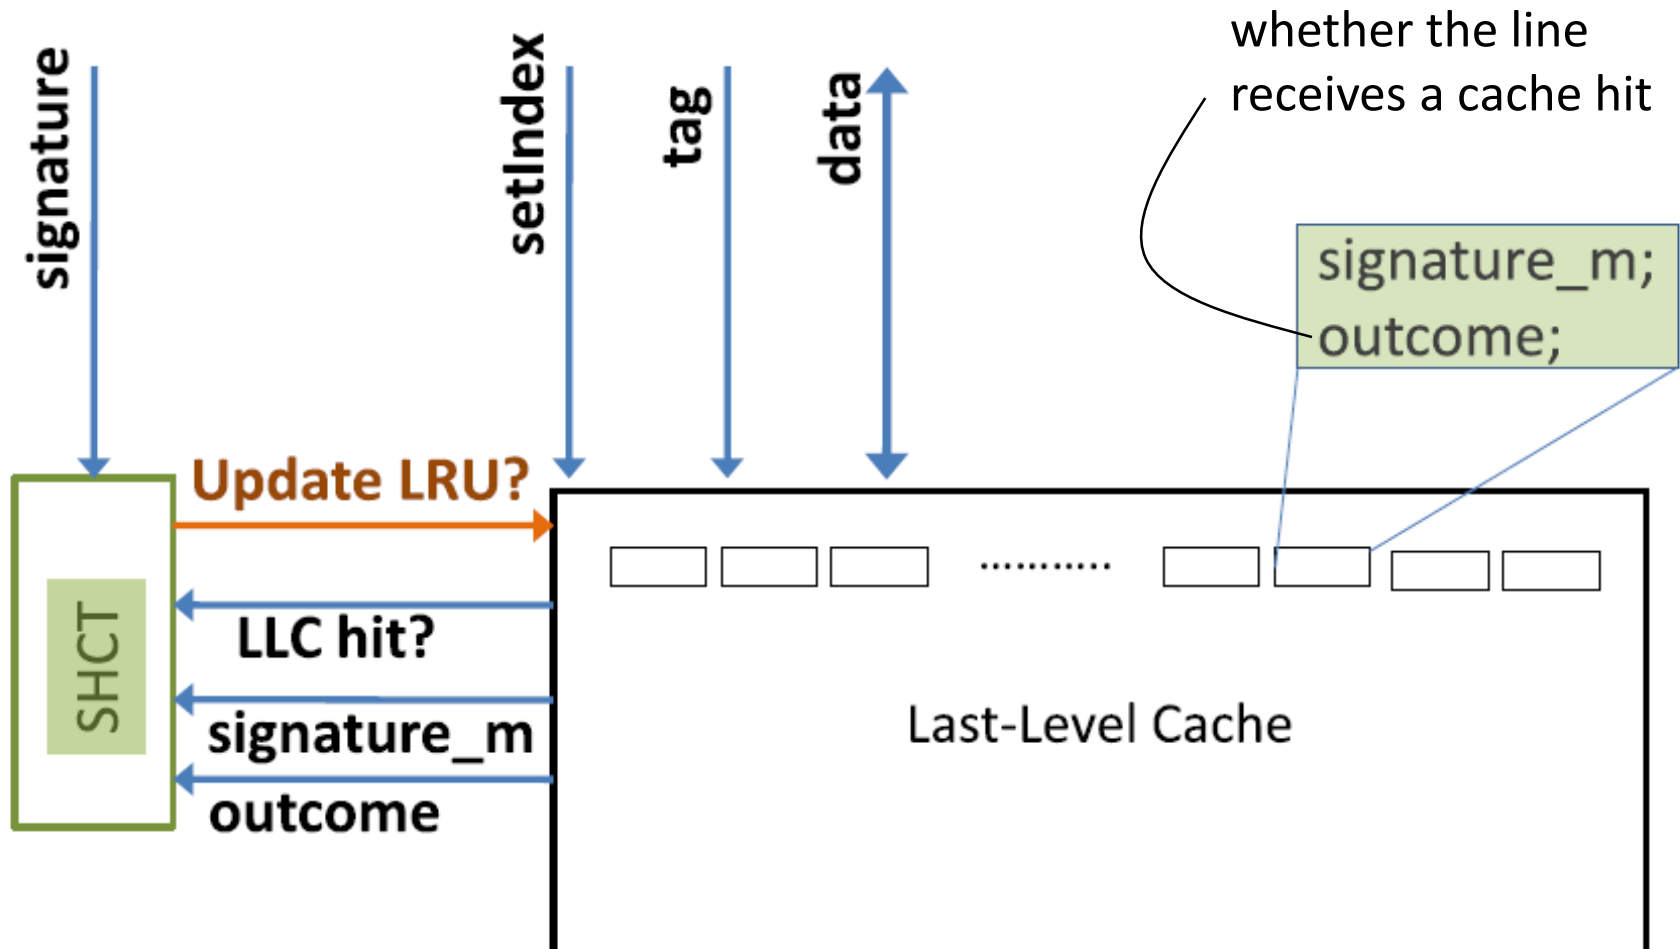
    // Update the stack position of all lines before the current line
    // Update implies incrementing their stack positions by one
    for(UINT32 way=0; way<assoc; way++)
    {
        if( repl[setIndex][way].LRUstackposition < currLRUstackposition )
        {
            repl[setIndex][way].LRUstackposition++;
        }
    }

    // Set the LRU stack position of new line to be zero
    repl[ setIndex ][ updateWayID ].LRUstackposition = 0;
}
```

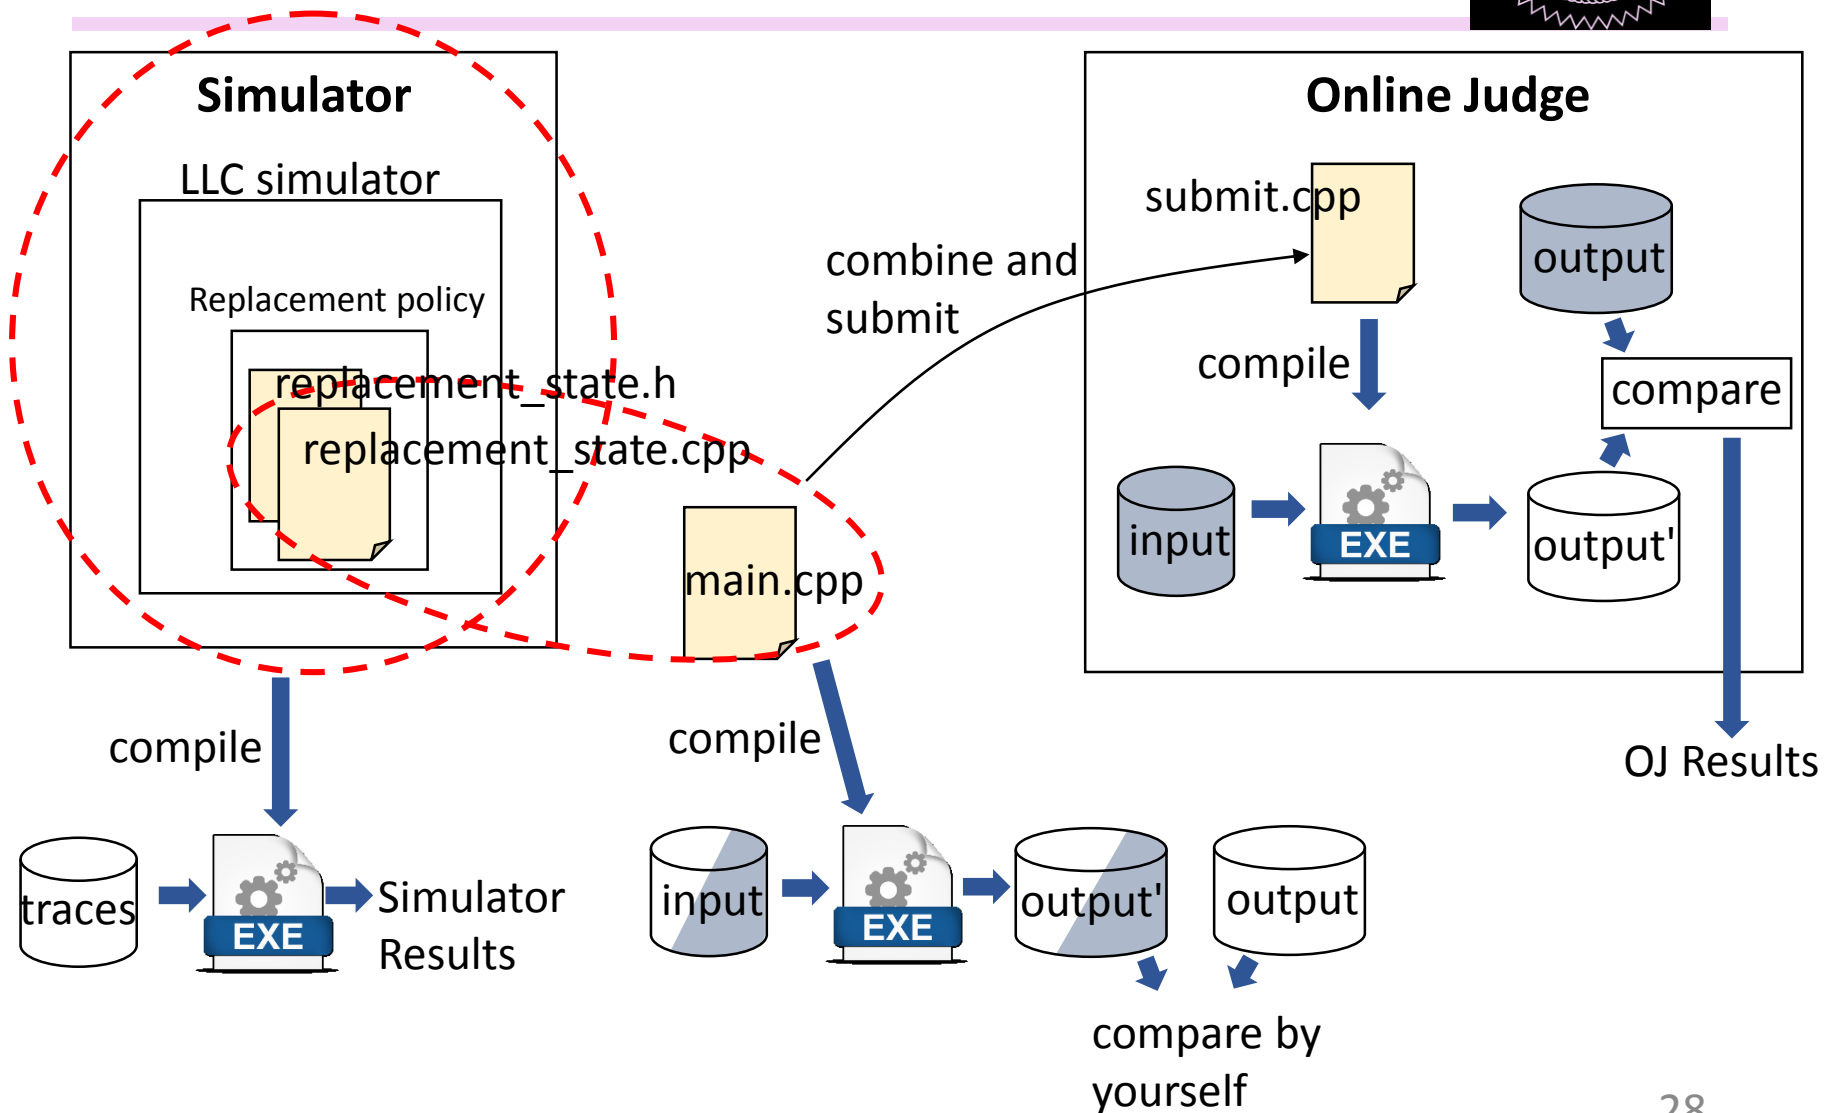
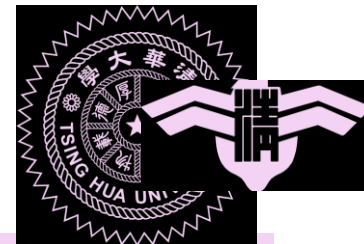
LRU Example



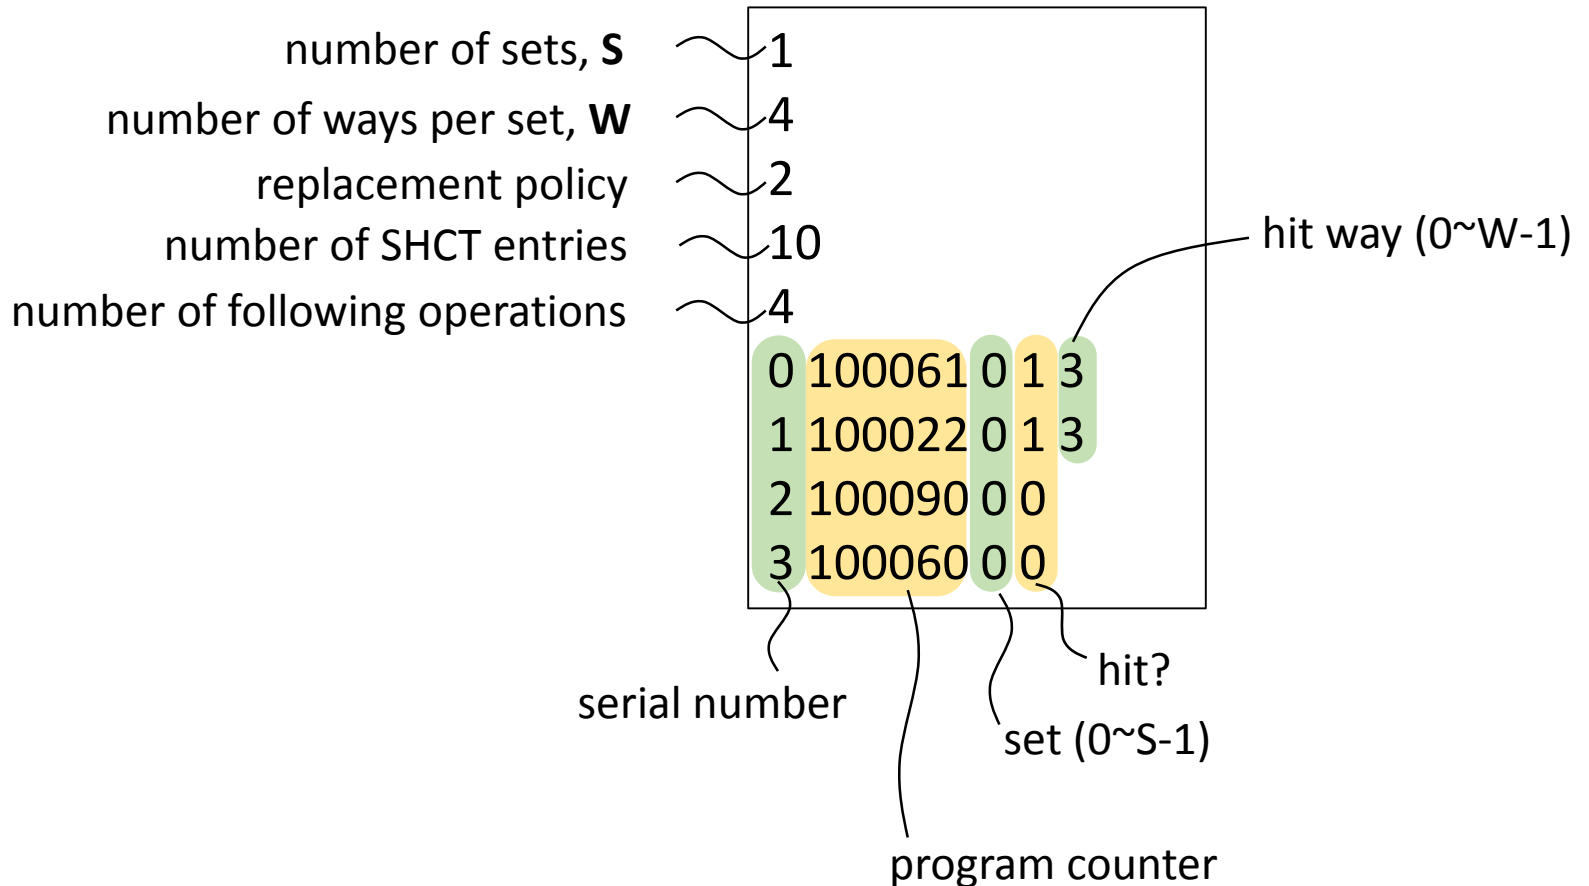
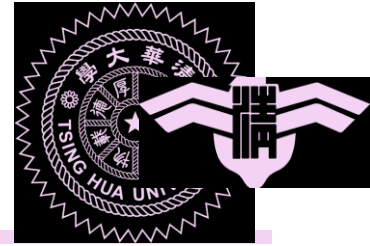
Implement SHiP



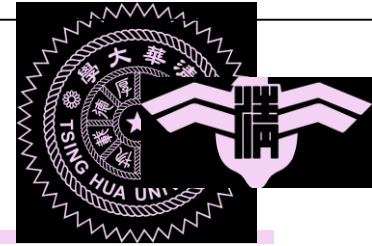
NTHU Online Judge



Input



Output



number of sets 1
number of ways per set 4
replacement policy 2

number of SHCT entries 10

number of following operations 4

serial number, program counter, hit/miss, hit way

0 100061 0 Hit 3

RRPV: 3 3 3 (0)

Signature: 0 0 0 0

outcome: 0 0 0 (1)

SHCT: (2) 1 1 1 1 1 1 1 1 1 1 ↵

↵

1 100022 0 Hit 3

RRPV: 3 3 3 (0)

Signature: 0 0 0 0

outcome: 0 0 0 (1)

SHCT: (3) 1 1 1 1 1 1 1 1 1 1

detailed status of the set

detailed status of the SHCT table

parenthesis indicates touched
value (not necessary changed)

Replaced way is shown for a miss

2 100090 0 Replace 0

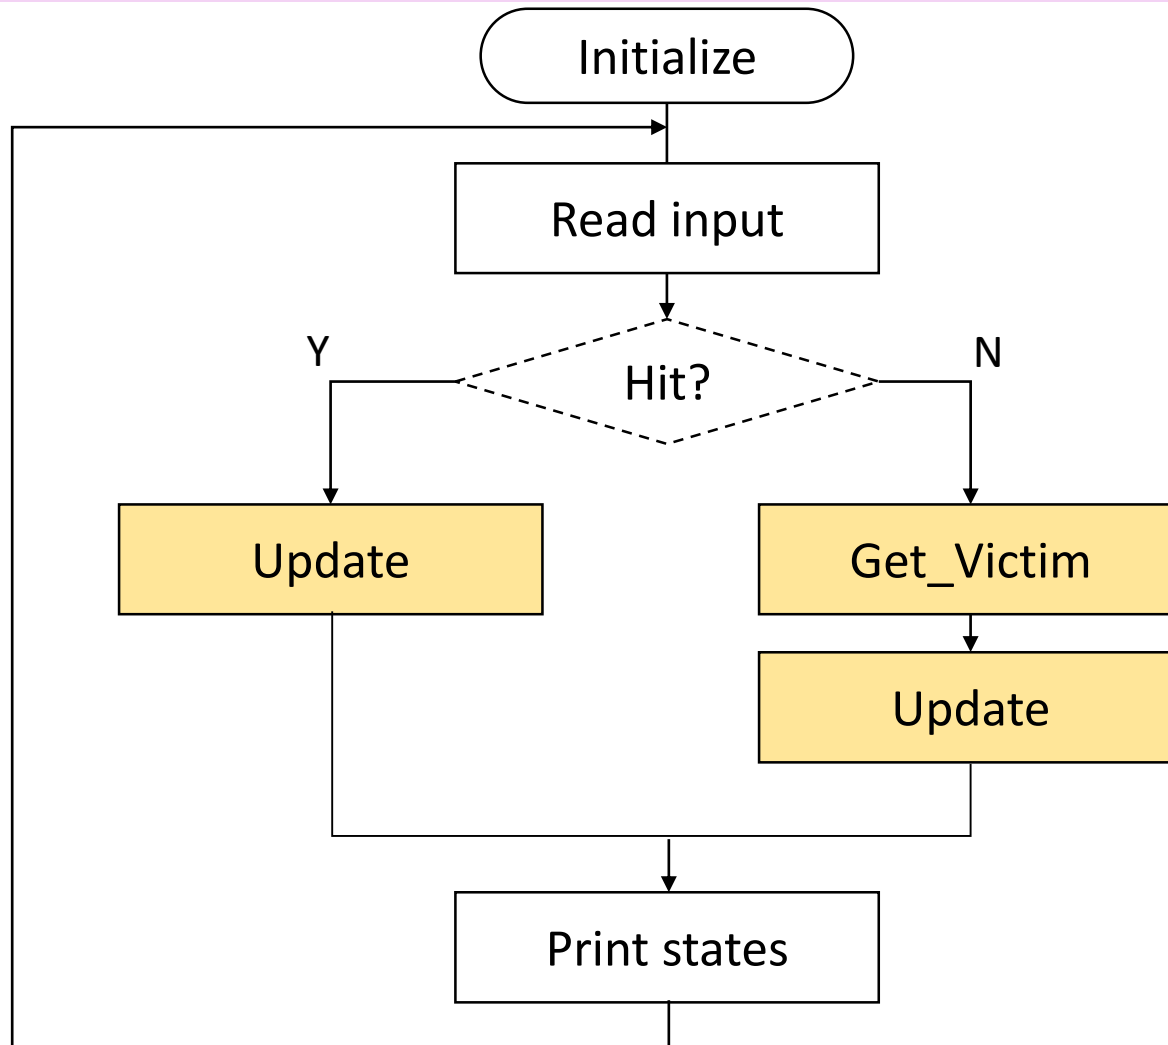
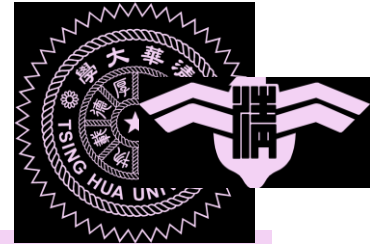
RRPV: (2) 3 3 0

Signature: (0) 0 0 0

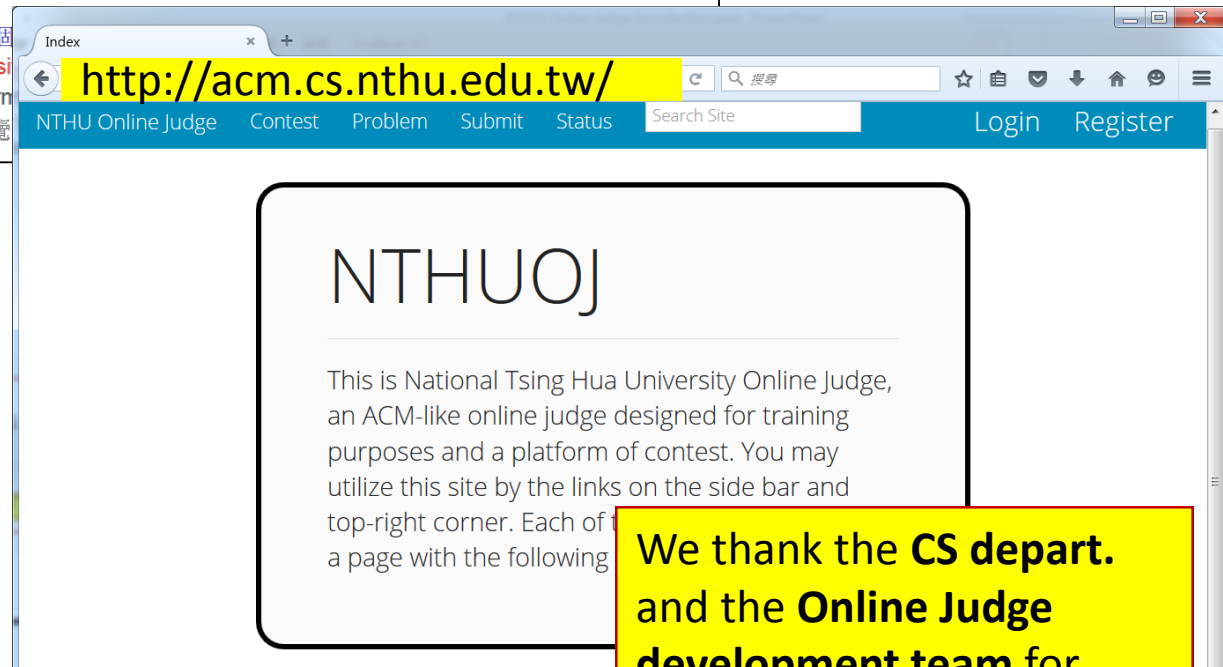
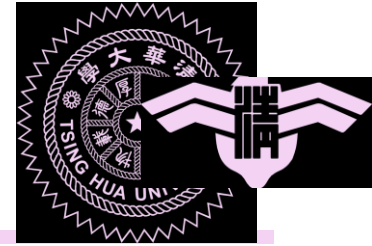
outcome: (0) 0 0 1

SHCT: (2) 1 1 1 1 1 1 1 1 1 1 ↵

main.cpp

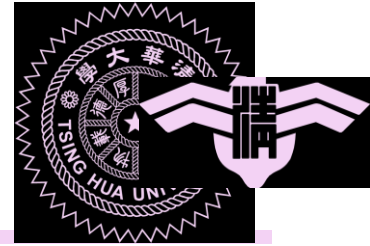


NTHU Online Judge

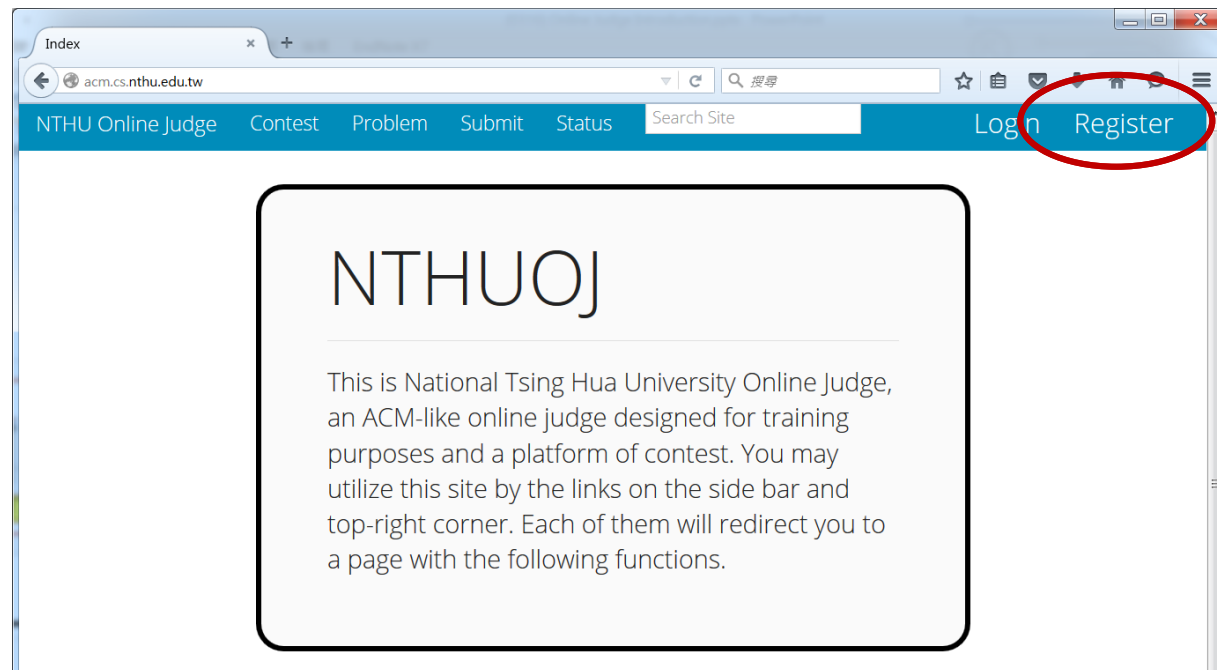


**We thank the CS depart.
and the Online Judge
development team for
supporting this service.**

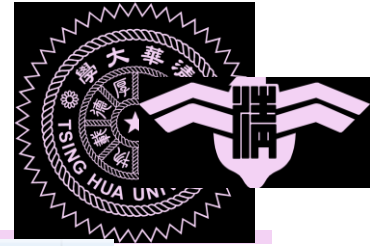
Login



- Register your ID:
MS+studentID
e.g., **MS**10203040506



Homework Page



<http://acm.cs.nthu.edu.tw/problem/10967/>

NTHU Online Judge

10967 - SHiP Cache

Status

Limits

Submit

Description

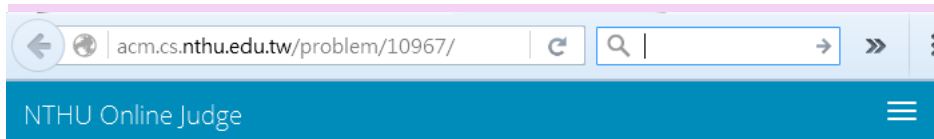
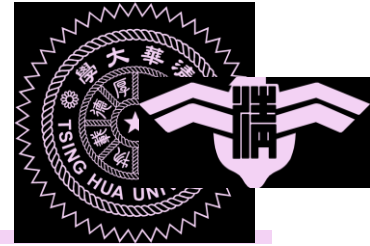
Simulating the SHiP (Signature-based Hit Prediction) policy.

Input

The first five lines are as follows:

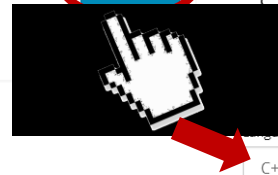
1. The number of sets of the cache.
2. The number of ways of the cache.
3. "2", which indicates the SHiP policy.
4. The number of saturating counters used by the SHiP policy.
5. The number of operations to simulate.

Submit Your Code



10967 - SHiP Cache

Status | Limits | **Submit**



Description

Simulating the SHiP (Signature-based Hit Prediction) policy.

Input

The first five lines are as follows:

1. The number of sets of the cache.
2. The number of ways of the cache.
3. "2", which indicates the SHiP policy.
4. The number of saturating counters used by the SHiP policy.
5. The number of operations to simulate.



Submit Code 10967 - SHiP Cache

Problem ID: 10967

Language: C++: -O2 -lm -std=c++

Backend:
gcc:
g++:

Code

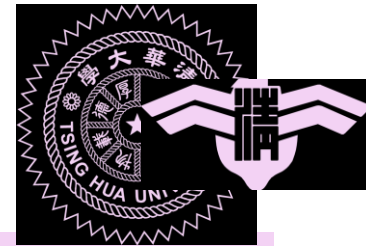
1 Code goes here...

Paste your code here

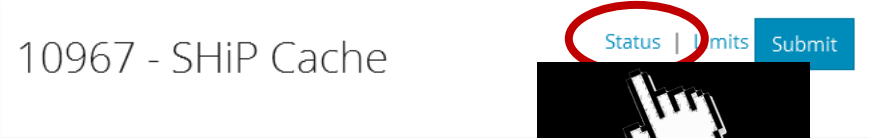
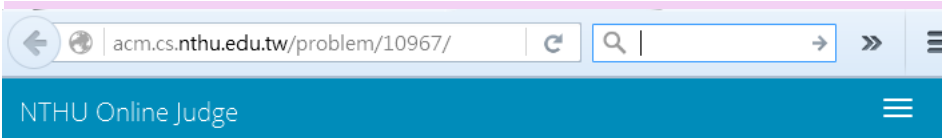
We support Sublime key bindings!

Browse





Check Your Results



Description

Simulating the SHiP (Signature-based Hit Prediction) policy.

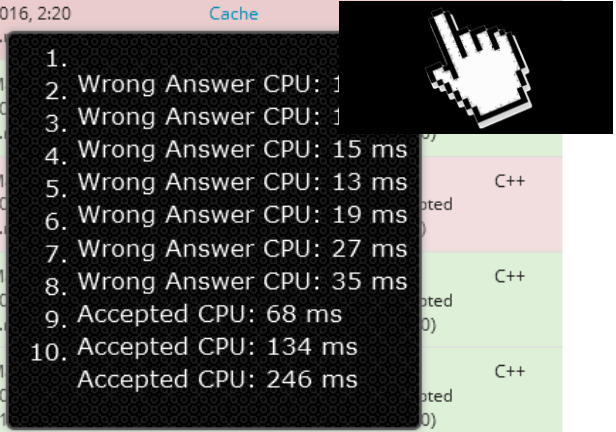
Input

The first five lines are as follows:

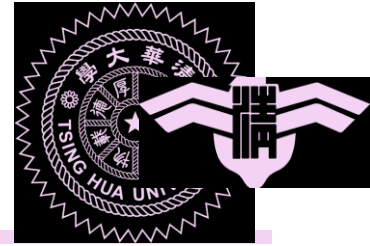
1. The number of sets of the cache.
2. The number of ways of the cache.
3. "2", which indicates the SHiP policy.
4. The number of saturating counters used by the SHiP policy.
5. The number of operations to simulate.



SID	Submit Time	Username	Problem	Status	Source
1347751	March 16, 2016, 4:01 p.m.	judgeDS2	10967 - SHiP Cache	All Accepted (10/10)	C++
1347738	March 16, 2016, 2:20 p.m.	judgeDS2	10967 - SHiP Cache	Not Accepted	C++
1347736	March 16, 2016, 2:20 p.m.	judgeDS2	10967 - SHiP Cache	Wrong Answer CPU: 15 ms	C++
1347732	March 16, 2016, 2:20 p.m.	judgeDS2	10967 - SHiP Cache	Wrong Answer CPU: 13 ms	C++
1347731	March 16, 2016, 2:20 p.m.	judgeDS2	10967 - SHiP Cache	Wrong Answer CPU: 19 ms	C++
1347721	March 16, 2016, 2:20 p.m.	judgeDS2	10967 - SHiP Cache	Wrong Answer CPU: 27 ms	C++
1347720	March 16, 2016, 11:44 a.m.	judgeDS2	10967 - SHiP Cache	Wrong Answer CPU: 35 ms	C++
1347649	March 15, 2016, 6:17 p.m.	judgeDS2	10967 - SHiP Cache	Accepted CPU: 68 ms	C++
1347648	March 15, 2016, 6:03 p.m.	judgeDS2	10967 - SHiP Cache	Accepted CPU: 134 ms	C++
1347647	March 15, 2016, 6:03 p.m.	judgeDS2	10967 - SHiP Cache	Accepted CPU: 246 ms	C++

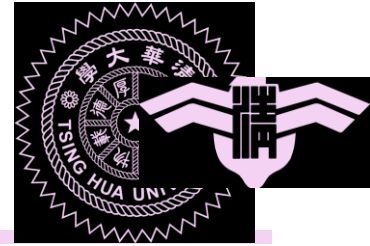


Hints



- Recommend you to develop, test, and debug programs in the Linux environment
- Online Judge can display error types (wrong answer, compiler error, etc.) and compiler errors
- But Online Judge does not show the exact error situation (e.g., the inputs, expected outputs, and your outputs)

Hints



- Use tools to help comparing your answer with the expected answer
- Non-standard libraries and some functions such as file IO are restricted in NTHU OJ. Using them leads to a "restricted function" error
- Extra space, newline characters, etc. lead to "representation error"
 - Online judge performs character to character matching to determine whether the output is correct



Grading Policy

- Grading
 - 100% Online judge results (including submit.cpp)
 - Please make sure that I can perform simulation using the original simulator code and your code
 - -10% per week till 70%



Other Details

- Other details will be announced at eeclass
- Please post your questions in the eeclass 討論區
- Please do not share or post your code or copy other students' code
 - Zero credit for all equivalent code